

# Facial Emotion Recognition

---

This repository is my project hand-in for the AKT3 course on Deep Learning & Computer Vision.

## Dataset


For training this model we will be using the [FER2013](#) dataset.

## Example Data

The dataset contains 48x48 images of human faces.



## Analysis



Using the distribution we can determine a baseline accuracy.

$$\text{HappyCounts} / \text{TotalCounts} = 0.25$$

$$\text{Baseline accuracy} = 25\%$$

So by always guessing **Happy** we could reach an accuracy of 25%. Our goal is to improve that with the CNN.

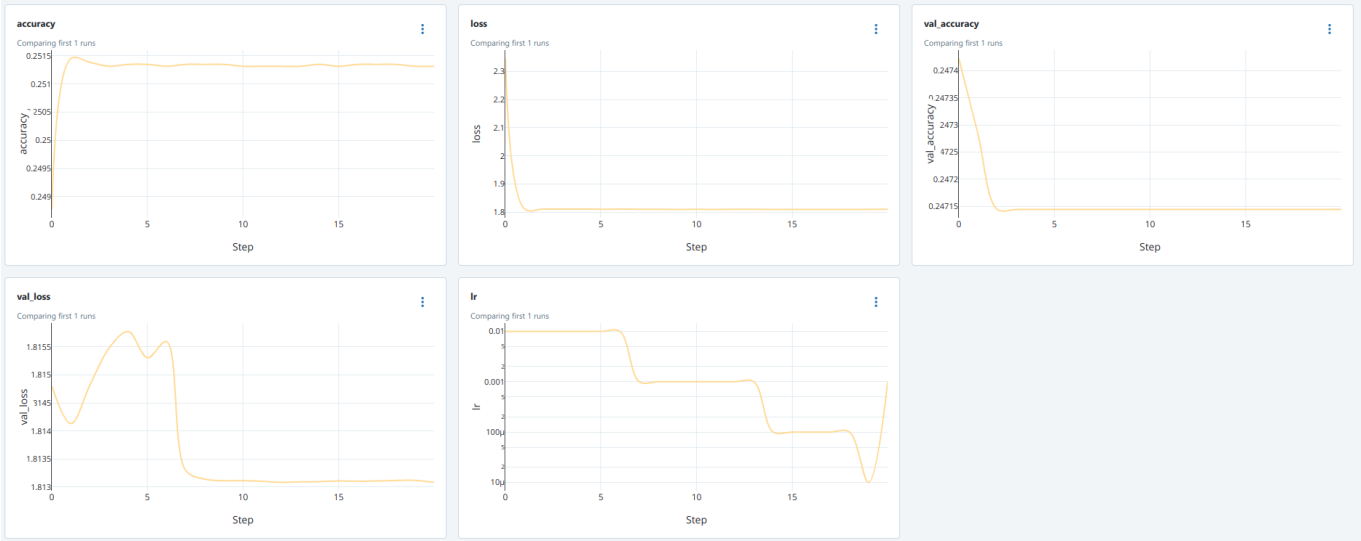
## Baseline

Using the `train.py` script we are training a Facial emotion Recognition model that classifies images of human faces on 7 emotions ("**Angry**", "**Disgust**", "**Fear**", "**Happy**", "**Sad**", "**Surprise**", "**Neutral**").

We split up the dataset into train, validation and test data.

- Train dataset size: 25120 examples
- Validation dataset size: 7179 examples
- Test dataset size: 3588 examples

## Results



As shown in the graphs above we achieve very poor performance with our baseline parameters.

Parameter	Value
learning_rate	0.01
loss	categorical_crossentropy
epochs	50
batch_size	128
early_stopping_patience	7
lr_patience	5
lr_reduction_factor	0.1
optimizer	Adam
num_classes	7
input_shape	(48, 48, 1)
shuffle	True
restore_best_weights	True


## Experiment 1 - Improving validation-accuracy

In my first run the model only achieved a validation accuracy of 21% which is very poor. I was confused because other resources showed me that on this dataset significantly higher validation accuracies with similar CNNs could be achieved.

My hypothesis is that I chose a far to high starting learning rate which lead to very early convergence and therefore significant underfitting. By reducing the learning rate I expect better results.

Parameter	Value
learning_rate	0.001

Results



As we can see in the resulting charts my hypothesis was correct and by reducing the learning rate we achieve much better results.


Experiment 2 - Smoothing the validation-loss curve

The new validation loss curve is very erratic. I want to make it smoother and reduce the bumpiness of the curve. For this I again will lower the learning rate by a factor of 10.

Parameter	Value
learning_rate	0.0001

Results

As we can see in the resulting charts my hypothesis was correct and by reducing the learning rate the curve is much less erratic.



Experiment 3 - Disabling the restore-best-weights option

For some reason in the EarlyStopping callback the restore\_best\_weights option actually chooses a worse configuration of the model at the end. By disabling the option we want to prohibit that behaviour.

```
keras.callbacks.EarlyStopping(  
    patience=params["early_stopping_patience"],  
    restore_best_weights=False  
)
```

Parameter	Value
restore_best_weights	False

Results

In the following graph we see that by disabling the `restore_best_weights` option we can actually keep the better model in the end.

